

Frequent Pattern-Growth with MPI

Generated by Doxygen 1.8.17

1 File Index	1
1.1 File List	1
2 File Documentation	2
2.1 src/io.h File Reference	2
2.1.1 Detailed Description	3
2.1.2 Function Documentation	3
2.2 src/reduce.h File Reference	6
2.2.1 Detailed Description	7
2.2.2 Function Documentation	7
2.3 src/sort.h File Reference	15
2.3.1 Detailed Description	16
2.3.2 Macro Definition Documentation	16
2.3.3 Function Documentation	16
2.4 src/tree.h File Reference	19
2.4.1 Detailed Description	21
2.4.2 Data Structure Documentation	21
2.4.3 Macro Definition Documentation	21
2.4.4 Typedef Documentation	21
2.4.5 Function Documentation	22
2.5 src/types.h File Reference	26
2.5.1 Detailed Description	27
2.5.2 Typedef Documentation	27
2.6 src/utils.h File Reference	28
2.6.1 Detailed Description	28
2.6.2 Function Documentation	28

1 File Index

1.1 File List

Here is a list of all files with brief descriptions:

src/ io.h	
Functions that handle the input and output of the program	2
src/ reduce.h	
Functions that perform the exchange of data between MPI processes	6

src/sort.h	
Functions that implement parallel QuickSort with OpenMP	15
src/tree.h	
Definition of FP-Tree and functions to build it	19
src/types.h	
Definition of datatypes used by the program	26
src/utils.h	
Utility functions	28

2 File Documentation

2.1 src/io.h File Reference

Functions that handle the input and output of the program.

```
#include "types.h"
#include <mpi.h>
```

Functions

- void [update_supports](#) ([Item](#) item, [SupportMap](#) *support_map)
Increase the support of the given item in the map. If the item is not present it is inserted with support 1.
- void [transaction_free](#) ([Transaction](#) *transaction)
Free the given transaction and all the items in it.
- void [transactions_free](#) ([TransactionsList](#) *transactions)
Free all the transaction in the list and the list itself.
- void [transactions_write](#) (int rank, [TransactionsList](#) transactions)
Write a list of transactions to the file named as the rank of the process.
- int [item_parse](#) (int rank, int i, char *chunk, int chunk_size, [Transaction](#) *transaction, [SupportMap](#) *support_map)
Parse an item from the string chunk, starting from position i up to the first space, newline or '\0'. The item is added to the transaction and the corresponding support is increased.
- int [transaction_parse](#) (int rank, int i, char *chunk, int chunk_size, [TransactionsList](#) *transactions, [SupportMap](#) *support_map)
Parse an transaction from the string chunk, starting from position i up to the first newline or '\0'. The transaction is added to the list of transactions. The support of the items in the transaction is increased as they get read.
- void [read_chunk](#) (char *filename, int rank, int world_size, char **chunk, int *my_size, int *read_size)
Read a chunk of the given file.
- void [transactions_read](#) ([TransactionsList](#) *transactions, char *filename, int rank, int world_size, [SupportMap](#) *support_map)
Read a list of transactions from the portion of file assigned to the current process. The support of the items read is increased in the support_map.

2.1.1 Detailed Description

Functions that handle the input and output of the program.

2.1.2 Function Documentation

2.1.2.1 item_parse() `int item_parse (`
 `int rank,`
 `int i,`
 `char * chunk,`
 `int chunk_size,`
 `Transaction * transaction,`
 `SupportMap * support_map)`

Parse an item from the string chunk, starting from position i up to the first space, newline or '\0'. The item is added to the transaction and the corresponding support is increased.

Parameters

<i>rank</i>	Rank of the current process
<i>i</i>	Start position from where to start parsing
<i>chunk</i>	String containing the item to parse
<i>chunk_size</i>	Size of the chunk
<i>transaction</i>	The transaction where to add the item
<i>support_map</i>	The map from items to respective support

Returns

The index where the parsed item ends (excluded)

2.1.2.2 read_chunk() `void read_chunk (`
 `char * filename,`
 `int rank,`
 `int world_size,`
 `char ** chunk,`
 `int * my_size,`
 `int * read_size)`

Read a chunk of the given file.

Read up to 2 * my_size, since we do not know where transactions start exactly.

Parameters

<i>filename</i>	File where transactions are stored
<i>rank</i>	Rank of the current process
<i>world_size</i>	Number of active processes
<i>chunk</i>	String where to store the read bytes
<i>my_size</i>	Number of bytes each process is assigned to
<i>read_size</i>	Number of bytes the current process has read

2.1.2.3 transaction_free() `void transaction_free (`
`Transaction * transaction)`

Free the given transaction and all the items in it.

Parameters

<i>transaction</i>	The transaction to free
--------------------	-------------------------

2.1.2.4 transaction_parse() `int transaction_parse (`
`int rank,`
`int i,`
`char * chunk,`
`int chunk_size,`
`TransactionsList * transactions,`
`SupportMap * support_map)`

Parse an transaction from the string chunk, starting from position i up to the first newline or '\0'. The transaction is added to the list of transactions. The support of the items in the transaction is increased as they get read.

Parameters

<i>rank</i>	Rank of the current process
<i>i</i>	Start position from where to start parsing
<i>chunk</i>	String containing the item to parse
<i>chunk_size</i>	Size of the chunk
<i>transactions</i>	The list of transactions where to add the transaction
<i>support_map</i>	The map from items to respective support

Returns

The index where the parsed transaction ends (excluded)

2.1.2.5 transactions_free() void transactions_free (
 TransactionsList * transactions)

Free all the transaction in the list and the list itself.

Parameters

<i>transactions</i>	The list of transactions to free
---------------------	----------------------------------

2.1.2.6 transactions_read() void transactions_read (
 TransactionsList * transactions,
 char * filename,
 int rank,
 int world_size,
 SupportMap * support_map)

Read a list of transactions from the portion of file assigned to the current process. The support of the items read is increased in the support_map.

Parameters

<i>transactions</i>	List of transactions where to store the data
<i>filename</i>	Name of the file from which to read
<i>rank</i>	Rank of the current process
<i>world_size</i>	Number of active processes
<i>support_map</i>	A map from items to the respective support

2.1.2.7 transactions_write() void transactions_write (
 int rank,
 TransactionsList transactions)

Write a list of transactions to the file named as the rank of the process.

Parameters

<i>rank</i>	Rank of the current process
<i>transactions</i>	List of transactions to write

2.1.2.8 update_supports() void update_supports (
 Item item,
 SupportMap * support_map)

Increase the support of the given item in the map. If the item is not present it is inserted with support 1.

Parameters

<i>item</i>	Item of wich to increase the support
<i>support_map</i>	A map from items to the respective support

2.2 src/reduce.h File Reference

Functions that perform the exchange of data between MPI processes.

```
#include "mpi.h"
#include "tree.h"
#include "types.h"
```

Functions

- MPI_Datatype [define_datatype_hashmap_element](#) ()
Define a datatype for an hashmap element in order to be able to send it with MPI.
- MPI_Datatype [define_datatype_tree_node](#) ()
Define an MPI Datatype for a [TreeNode](#) in order to be able to send it with MPI.
- void [merge_map](#) (SupportMap *support_map, hashmap_element *elements, int size)
Merge the current map with an array of elements.
- void [recv_map](#) (int rank, int world_size, int source, SupportMap *support_map, MPI_↔
 Datatype DT_HASHMAP_ELEMENT)
Receive an array of hashmap elements with MPI.
- void [send_map](#) (int rank, int world_size, int dest, SupportMap *support_map, MPI_↔
 Datatype DT_HASHMAP_ELEMENT)
Send an array of hashmap elements with MPI.

- void [broadcast_map](#) (int rank, int world_size, [SupportMap](#) *support_map, hashmap_element **items_count, int *num_items, MPI_Datatype DT_HASHMAP_ELEMENT, int min_support)
Broadcast all the elements of a SupportMap to every MPI process.
- void [get_global_map](#) (int rank, int world_size, [SupportMap](#) *support_map, hashmap_element **items_count, int *num_items, int min_support)
Get the global map object for every MPI process.
- void [merge_indices](#) (int rank, int *sorted_indices, int start1, int end1, int start2, int end2, hashmap_element *items_count, int num_items)
Merge two arrays of sorted indices into a unique array of sorted indices.
- void [recv_indices](#) (int rank, int source, int *sorted_indices, int start, int *end, int length, int size, hashmap_element *items_count, int num_items)
Receive an array of sorted indices.
- void [send_indices](#) (int rank, int dest, int *sorted_indices, int start, int end, int length)
Send an array of sorted indices.
- void [broadcast_indices](#) (int rank, int *sorted_indices, int num_items)
Broadcast the final global array of indices to every MPI process in the world.
- void [get_sorted_indices](#) (int rank, int world_size, int *sorted_indices, int start, int end, int length, hashmap_element *items_count, int num_items)
Get the global sorted indices array for every MPI process.
- void [parse_tree](#) ([TreeNodeToSend](#) *nodes, int num_nodes, [Tree](#) *dest)
Parse an array of TreeNodesToSend into a Tree structure.
- void [send_tree](#) (int dest, [Tree](#) *tree, MPI_Datatype DT_TREE_NODE)
Sends a tree to an MPI process and frees up the memory.
- void [recv_tree](#) (int source, [Tree](#) *tree, MPI_Datatype DT_TREE_NODE)
Receive a tree from an MPI process.
- void [broadcast_tree](#) (int rank, [Tree](#) *tree, MPI_Datatype DT_TREE_NODE)
Broadcast the final FP-Tree to every MPI process in the world.
- void [get_global_tree](#) (int rank, int world_size, [Tree](#) *tree)
Get the global FP-tree on every MPI process.

2.2.1 Detailed Description

Functions that perform the exchange of data between MPI processes.

2.2.2 Function Documentation

2.2.2.1 broadcast_indices() void broadcast_indices (
 int rank,
 int * sorted_indices,
 int num_items)

Broadcast the final global array of indices to every MPI process in the world.

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>sorted_indices</i>	The array of sorted indices that has to be broadcasted
<i>num_items</i>	The number of elements of the sorted_indices array

2.2.2.2 broadcast_map() void broadcast_map (

```

    int rank,
    int world_size,
    SupportMap * support_map,
    hashmap_element ** items_count,
    int * num_items,
    MPI_Datatype DT_HASHMAP_ELEMENT,
    int min_support )

```

Broadcast all the elements of a SupportMap to every MPI process.

This is used by the master process to send the global version of the map to every MPI process in the world

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>world_size</i>	The number of MPI processes in the world
<i>support_map</i>	The map from which to extract the elements to be sent
<i>items_count</i>	A pointer to an array of hashmap elements having the item string as a key and the support count as a value
<i>num_items</i>	A pointer to an integer describing the total number of items in the broadcasted map
<i>DT_HASHMAP_ELEMENT</i>	An MPI datatype that describes the data structure that has to be sent with MPI
<i>min_support</i>	The minimum support that an item has to have in order to be contained in the final map

2.2.2.3 broadcast_tree() void broadcast_tree (

```

    int rank,
    Tree * tree,
    MPI_Datatype DT_TREE_NODE )

```

Broadcast the final FP-Tree to every MPI process in the world.

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>tree</i>	The tree that has to be broadcasted in case the process rank is 0, otherwise the tree which the final tree will be stored
<i>DT_TREE_NODE</i>	MPI_Datatype describing a TreeNode

2.2.2.4 `define_datatype_hashmap_element()` `MPI_Datatype define_datatype_hashmap↵_element ()`

Define a datatype for an hashmap element in order to be able to send it with MPI.

Returns

MPI_Datatype describing an hashmap element

2.2.2.5 `define_datatype_tree_node()` `MPI_Datatype define_datatype_tree_node ()`

Define an MPI Datatype for a [TreeNode](#) in order to be able to send it with MPI.

Returns

MPI_Datatype describing a [TreeNode](#)

2.2.2.6 `get_global_map()` `void get_global_map (` `int rank,` `int world_size,` `SupportMap * support_map,` `hashmap_element ** items_count,` `int * num_items,` `int min_support)`

Get the global map object for every MPI process.

This function follows a tree-like structure to build the global map of items of our program. Initially, every process has its own partial map of items. Then, processes which meet the condition ($\text{rank} \% \text{pow} \neq 0$), where pow is a function of the current level in the tree, sends its partial map to the specified destination. On the contrary, processes which meet the condition ($\text{rank} \% \text{pow} == 0$) receive the map from the other processes, and then merge the received elements with their current map. This is done until the only process in the current level is the process number 0. Upon reaching that state, it means that process 0 now has the complete map of every process and can broadcast its knowledge to the whole domain.

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>world_size</i>	The number of MPI processes in the world
<i>support_map</i>	The map from which to extract the elements to be sent
<i>items_count</i>	A pointer to an array of hashmap elements having the item string as a key and the support count as a value
<i>num_items</i>	A pointer to an integer describing the total number of items in the broadcasted map
<i>min_support</i>	The minimum support that an item has to have in order to be contained in the final map

2.2.2.7 get_global_tree() `void get_global_tree (`
`int rank,`
`int world_size,`
`Tree * tree)`

Get the global FP-tree on every MPI process.

This function follows a tree-like structure to build the final FP-tree. Initially, every process has its own partial tree composed by the transactions of its local assigned chunk. Then, processes which meet the condition ($\text{rank} \% \text{pow} \neq 0$), where *pow* is a function of the current level in the tree, sends its partial tree to the specified destination. On the contrary, processes which meet the condition ($\text{rank} \% \text{pow} == 0$) receive the FP-tree from the other processes, and then merge the received partial FP-tree with their current FP-tree. This is done until the only process in the current level is the process number 0. Upon reaching that state, it means that process 0 now has the complete FP-tree and can broadcast its knowledge to the whole domain.

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>world_size</i>	The number of processes in the current world
<i>tree</i>	The tree that has to be sent/received. This structure is heavily manipulated during the execution of this function.

2.2.2.8 get_sorted_indices() `void get_sorted_indices (`
`int rank,`
`int world_size,`
`int * sorted_indices,`
`int start,`
`int end,`

```

int length,
hashmap_element * items_count,
int num_items )

```

Get the global sorted indices array for every MPI process.

This function follows a tree-like structure to build the global sorted array of indices of the items of our program. Initially, every process has its own partial array of sorted indices. Then, processes which meet the condition ($\text{rank} \% \text{pow} \neq 0$), where pow is a function of the current level in the tree, sends its partial array to the specified destination. On the contrary, processes which meet the condition ($\text{rank} \% \text{pow} == 0$) receive the array from the other processes, and then merge the received elements with their current array. This is done until the only process in the current level is the process number 0. Upon reaching that state, it means that process 0 now has the complete array of every process and can broadcast its knowledge to the whole domain.

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>world_size</i>	The number of MPI processes in the current world
<i>sorted_indices</i>	The array of sorted indices that will be used to send the data
<i>start</i>	The starting index of the items that the function will send
<i>end</i>	The ending index of the items that the function will send
<i>length</i>	The length of each process' original interval width
<i>items_count</i>	An array of hashmap elements having the item string as a key and the support count as a value
<i>num_items</i>	The number of items in the sorted_indices array

2.2.2.9 merge_indices() void merge_indices (

```

int rank,
int * sorted_indices,
int start1,
int end1,
int start2,
int end2,
hashmap_element * items_count,
int num_items )

```

Merge two arrays of sorted indices into a unique array of sorted indices.

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>sorted_indices</i>	The array of sorted indices that contains both arrays that have to be sorted
<i>start1</i>	The starting index of the first array
<i>end1</i>	The ending index of the first array

Parameters

<i>start2</i>	The starting index of the second array
<i>end2</i>	The ending index of the second array
<i>items_count</i>	An array of hashmap elements having the item string as a key and the support count as a value
<i>num_items</i>	The number of items in the sorted_indices array

2.2.2.10 merge_map() void merge_map (
 SupportMap * support_map,
 hashmap_element * elements,
 int size)

Merge the current map with an array of elements.

Parameters

<i>support_map</i>	A pointer to the map which will be populated with the merged items
<i>elements</i>	The array of hashmap elements to merge
<i>size</i>	The size of the elements array

2.2.2.11 parse_tree() void parse_tree (
 TreeNodeToSend * nodes,
 int num_nodes,
 Tree * dest)

Parse an array of TreeNodesToSend into a Tree structure.

Parameters

<i>nodes</i>	The array of TreeNodesToSend that have to be parsed
<i>num_nodes</i>	The size of the nodes array
<i>dest</i>	A pointer to the tree that will be created

2.2.2.12 recv_indices() void recv_indices (
 int rank,
 int source,

```

    int * sorted_indices,
    int start,
    int * end,
    int length,
    int size,
    hashmap_element * items_count,
    int num_items )

```

Receive an array of sorted indices.

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>source</i>	The rank of the MPI process that sends the data
<i>sorted_indices</i>	The array of sorted indices where the received data will be stored
<i>start</i>	The starting index of the items that the function will receive
<i>end</i>	The ending index of the items that the function will receive
<i>length</i>	The length of each process' original interval width
<i>size</i>	The current size of the received data
<i>items_count</i>	An array of hashmap elements having the item string as a key and the support count as a value
<i>num_items</i>	The number of items in the sorted_indices array

2.2.2.13 recv_map() void recv_map (

```

    int rank,
    int world_size,
    int source,
    SupportMap * support_map,
    MPI_Datatype DT_HASHMAP_ELEMENT )

```

Receive an array of hashmap elements with MPI.

First, receive the size of the array of hashmap elements and then receive the actual elements

Parameters

<i>rank</i>	MPI process rank
<i>world_size</i>	Number of MPI processes in the current world
<i>source</i>	The rank of the MPI process that sends the array of elements
<i>support_map</i>	A pointer to the map where the received elements have to be inserted
<i>DT_HASHMAP_ELEMENT</i>	An MPI datatype that describes the data structure received

2.2.2.14 recv_tree() void recv_tree (

```

    int source,
    Tree * tree,
    MPI_Datatype DT_TREE_NODE )

```

Receive a tree from an MPI process.

Parameters

<i>source</i>	The MPI process that is sending the data
<i>tree</i>	A pointer to partial tree of the current process, which will be integrated by merging the received tree
<i>DT_TREE_NODE</i>	MPI_Datatype describing a TreeNode

2.2.2.15 send_indices() void send_indices (

```

    int rank,
    int dest,
    int * sorted_indices,
    int start,
    int end,
    int length )

```

Send an array of sorted indices.

Parameters

<i>rank</i>	The rank of the process that broadcasts the data
<i>dest</i>	The rank of the MPI process that will receive the data
<i>sorted_indices</i>	The array of sorted indices that will be used to send the data
<i>start</i>	The starting index of the items that the function will send
<i>end</i>	The ending index of the items that the function will send
<i>length</i>	The length of each process' original interval width

2.2.2.16 send_map() void send_map (

```

    int rank,
    int world_size,
    int dest,
    SupportMap * support_map,
    MPI_Datatype DT_HASHMAP_ELEMENT )

```

Send an array of hashmap elements with MPI.

First, send the size of the array of hashmap elements and then send the actual elements.

Parameters

<i>rank</i>	MPI process rank
<i>world_size</i>	Number of MPI processes in the current world
<i>dest</i>	The rank of the MPI process that will receive the array of elements
<i>support_map</i>	A pointer to the map where the elements to be sent are stored
<i>DT_HASHMAP_ELEMENT</i>	An MPI datatype that describes the data structure that has to be sent with MPI

2.2.2.17 send_tree() void send_tree (
 int dest,
 Tree * tree,
 MPI_Datatype DT_TREE_NODE)

Sends a tree to an MPI process and frees up the memory.

Parameters

<i>dest</i>	The destination process that will receive the tree
<i>tree</i>	The tree that has to be sent
<i>DT_TREE_NODE</i>	MPI_Datatype describing a TreeNode

2.3 src/sort.h File Reference

Functions that implement parallel QuickSort with OpenMP.

```
#include "hashmap/hashmap.h"
#include "types.h"
```

Macros

- #define [INSERTION_SORT_THRESH](#) 100

Functions

- void [insertion_sort](#) (hashmap_element *items_count, int num_items, int *sorted_indices, int start, int end)

Sort the indices contained in the array sorted_indices from start to end using insertion sort. The indices are sorted according to the corresponding element in items_count.

- void `swap` (int *a, int i, int j)
Swap elements in position i and j in the array a.
- int `pivot` (hashmap_element *items_count, int num_items, int *sorted_indices, int start, int end, int m)
Perform the pivot operation of the quicksort algorithm on the sub-array sorted_indices from start to end (included).
- int `choose_pivot` (hashmap_element *items_count, int num_items, int *sorted_indices, int start, int end)
Choose a pivot according to the Median-of-three heuristic.
- void `parallel_quick_sort` (hashmap_element *items_count, int num_items, int *sorted_indices, int start, int end, int **stack, int *num_busy_threads, int *num_threads)
Parallel version of the QuickSort algorithm.
- void `sort` (hashmap_element *items_count, int num_items, int *sorted_indices, int start, int end, int num_threads)
Put in the array sorted_indices the indices of the elements of the array items_count from position start to end, after they are sorted according to their value field.

2.3.1 Detailed Description

Functions that implement parallel QuickSort with OpenMP.

2.3.2 Macro Definition Documentation

2.3.2.1 INSERTION_SORT_THRESH `#define INSERTION_SORT_THRESH 100`

2.3.3 Function Documentation

2.3.3.1 `choose_pivot()`

```
int choose_pivot (  
    hashmap_element * items_count,  
    int num_items,  
    int * sorted_indices,  
    int start,  
    int end )
```

Choose a pivot according to the Median-of-three heuristic.

The pivot is chosen as the median of the first, middle and last element of the sub-array sorted_indices from start to end, according to the corresponding value in items_count. These three indices are swapped, so that the median at the end is in position start.

Parameters

<i>items_count</i>	array of key-value pairs
<i>num_items</i>	length of <i>items_count</i>
<i>sorted_indices</i>	array of indices that are going to be sorted
<i>start</i>	start position of the sub-array to sort
<i>end</i>	end position of the sub-array to sort

Returns

the position of the pivot

2.3.3.2 insertion_sort() void insertion_sort (
hashmap_element * *items_count*,
int *num_items*,
int * *sorted_indices*,
int *start*,
int *end*)

Sort the indices contained in the array *sorted_indices* from *start* to *end* using insertion sort. The indices are sorted according to the corresponding element in *items_count*.

Parameters

<i>items_count</i>	array of key-value pairs
<i>num_items</i>	length of <i>items_count</i>
<i>sorted_indices</i>	array of indices that are going to be sorted
<i>start</i>	start position of the sub-array to sort
<i>end</i>	end position of the sub-array to sort

2.3.3.3 parallel_quick_sort() void parallel_quick_sort (
hashmap_element * *items_count*,
int *num_items*,
int * *sorted_indices*,
int *start*,
int *end*,
int ** *stack*,
int * *num_busy_threads*,
int * *num_threads*)

Parallel version of the QuickSort algorithm.

Parameters

<i>items_count</i>	array of key-value pairs
<i>num_items</i>	length of <i>items_count</i>
<i>sorted_indices</i>	array of indices that are going to be sorted
<i>start</i>	start position of the sub-array to sort
<i>end</i>	end position of the sub-array to sort
<i>stack</i>	stack containing the jobs that need to be done by free threads
<i>num_busy_threads</i>	Number of currently busy threads
<i>num_threads</i>	Number of threads that perform the sorting

```

2.3.3.4 pivot() int pivot (
    hashmap_element * items_count,
    int num_items,
    int * sorted_indices,
    int start,
    int end,
    int m )

```

Perform the pivot operation of the quicksort algorithm on the sub-array *sorted_indices* from *start* to *end* (included).

The indices which value in *items_count* is smaller than the value of the pivot are moved to the left of the pivot, the ones with a value greater or equal are moved to its right.

Parameters

<i>items_count</i>	array of key-value pairs
<i>num_items</i>	length of <i>items_count</i>
<i>sorted_indices</i>	array of indices that are going to be sorted
<i>start</i>	start position of the sub-array to sort
<i>end</i>	end position of the sub-array to sort
<i>m</i>	position of the pivot

Returns

the new position of the pivot

```

2.3.3.5 sort() void sort (
    hashmap_element * items_count,

```

```

int num_items,
int * sorted_indices,
int start,
int end,
int num_threads )

```

Put in the array `sorted_indices` the indices of the elements of the array `items_count` from position `start` to `end`, after they are sorted according to their value field.

The array `items_count` is not modified. The algorithm implement a parallel version of QuickSort, described in the paper Süß M, Leopold C. A user's experience with parallel sorting and OpenMP. In Proceedings of the Sixth European Workshop on OpenMP-EWOMP'04 2004 Oct 18 (pp. 23-38).

Parameters

<i>items_count</i>	array of key-value pairs
<i>num_items</i>	length of <code>items_count</code>
<i>sorted_indices</i>	array of indices that are going to be sorted
<i>start</i>	start position of the sub-array to sort
<i>end</i>	end position of the sub-array to sort
<i>num_threads</i>	Number of threads that perform the sorting

2.3.3.6 swap() `void swap (`
`int * a,`
`int i,`
`int j)`

Swap elements in position `i` and `j` in the array `a`.

Parameters

<i>a</i>	array where to swap the elements
<i>i</i>	position of the first element to swap
<i>j</i>	position of the second element to swap

2.4 src/tree.h File Reference

Definition of FP-Tree and functions to build it.

```
#include "types.h"
```

Data Structures

- struct [TreeNode](#)
A node of an FP-Tree. [More...](#)
- struct [TreeNodeToSend](#)
A node of an FP-Tree to send with MPI. [More...](#)

Macros

- #define [TREE_NODE_NULL](#) -1

Typedefs

- typedef [TreeNode](#) ** [Tree](#)
FP-Tree.

Functions

- [TreeNode](#) * [tree_node_new](#) (int key, int value, int parent)
Instantiate a new node of the tree.
- void [tree_node_free](#) ([TreeNode](#) *node)
Free a tree node.
- [Tree](#) [tree_new](#) ()
Instantiate a new tree.
- void [tree_free](#) ([Tree](#) *tree)
Free the tree.
- int [tree_add_node](#) ([Tree](#) *tree, [TreeNode](#) *node)
Add a node to the tree.
- void [tree_add_subtree](#) ([Tree](#) *dest, [Tree](#) source, int nd, int ns)
Add the subtree rooted in the ns(th) node of the tree source as a child of the nd(th) node of the tree dest. The source tree is modified, as the nodes are moved to the destination tree.
- void [tree_merge_dfs](#) ([Tree](#) *dest, [Tree](#) source, int nd, int ns)
Merge the subtree of dest rooted in node with id nd with the subtree of source rooted in ns and store the result in dest. Also the source tree is modified.
- void [tree_merge](#) ([Tree](#) *dest, [Tree](#) source)
Merge the trees dest and source and store the result in dest. The source tree is modified. It is a wrapper for.
- void [tree_get_nodes](#) ([Tree](#) tree, [TreeNodeToSend](#) **nodes)
Inserts into the vector nodes the nodes to send.
- void [tree_print](#) ([Tree](#) tree)
Print the tree.
- [Tree](#) [tree_build_from_transaction](#) (int rank, int world_size, [Transaction](#) *transaction, [IndexMap](#) index_map, [hashmap_element](#) *items_count, int num_items, int *sorted_↔ indices)
Build a tree given a transaction.
- [Tree](#) [tree_build_from_transactions](#) (int rank, int world_size, [TransactionsList](#) transactions, [IndexMap](#) index_map, [hashmap_element](#) *items_count, int num_items, int *sorted_↔ indices, int num_threads)
Build a tree given a list of transactions.

2.4.1 Detailed Description

Definition of FP-Tree and functions to build it.

2.4.2 Data Structure Documentation

2.4.2.1 struct **TreeNode** A node of an FP-Tree.

Data Fields

map_t	adj	Adjacency map of the node, where the values are the indices of the children of the node in the tree and the keys are the ids of the corresponding items.
int	key	Id of the item represented by the node
int	parent	Index of the parent of the node in the tree
int	value	Value of the item represented by the node

2.4.2.2 struct **TreeNodeToSend** A node of an FP-Tree to send with MPI.

Data Fields

int	key	Id of the item represented by the node
int	parent	Index of the parent of the node in the tree
int	value	Value of the item represented by the node

2.4.3 Macro Definition Documentation

2.4.3.1 **TREE_NODE_NULL** `#define TREE_NODE_NULL -1`

2.4.4 Typedef Documentation

2.4.4.1 Tree `typedef TreeNode* * Tree`

FP-Tree.

2.4.5 Function Documentation

2.4.5.1 `tree_add_node()` `int tree_add_node (` `Tree * tree,` `TreeNode * node)`

Add a node to the tree.

Parameters

<i>tree</i>	Pointer to the tree
<i>node</i>	Pointer to the node to add

Returns

The id of the node in the tree

2.4.5.2 `tree_add_subtree()` `void tree_add_subtree (` `Tree * dest,` `Tree source,` `int nd,` `int ns)`

Add the subtree rooted in the ns(th) node of the tree source as a child of the nd(th) node of the tree dest. The source tree is modified, as the nodes are moved to the destination tree.

Parameters

<i>dest</i>	Pointer to the destination tree
<i>source</i>	Pointer to the source tree
<i>nd</i>	Id of the node in the destination tree
<i>ns</i>	Id of the node in the source tree

2.4.5.3 `tree_build_from_transaction()` `Tree tree_build_from_transaction (`

```

    int rank,
    int world_size,
    Transaction * transaction,
    IndexMap index_map,
    hashmap_element * items_count,
    int num_items,
    int * sorted_indices )

```

Build a tree given a transaction.

Parameters

<i>rank</i>	The rank of the process
<i>world_size</i>	The number of processes in the world
<i>transaction</i>	The transaction
<i>index_map</i>	The map from item to the corresponding id
<i>items_count</i>	The array of hashmap elements having the item string as a key and the support count as a value
<i>num_items</i>	The number of items in the sorted_indices array
<i>sorted_indices</i>	The array of the sorted indices of the items

Returns

The built tree

2.4.5.4 tree_build_from_transactions() `Tree` tree_build_from_transactions (

```

    int rank,
    int world_size,
    TransactionsList transactions,
    IndexMap index_map,
    hashmap_element * items_count,
    int num_items,
    int * sorted_indices,
    int num_threads )

```

Build a tree given a list of transactions.

First, we build the trees for the single transactions. Then, we merge them in a binary-tree-like fashion.

Parameters

<i>rank</i>	The rank of the process
<i>world_size</i>	The number of processes in the world
<i>transactions</i>	

Parameters

<i>index_map</i>	The map from item to the corresponding id
<i>items_count</i>	The array of hashmap elements having the item string as a key and the support count as a value
<i>num_items</i>	The number of items in the sorted_indices array
<i>sorted_indices</i>	The array of the sorted indices of the items
<i>num_threads</i>	The number of threads requested to perform the building

Returns

The built tree

2.4.5.5 tree_free() `void tree_free (`
`Tree * tree)`

Free the tree.

Parameters

<i>tree</i>	Pointer to the tree to free
-------------	-----------------------------

2.4.5.6 tree_get_nodes() `void tree_get_nodes (`
`Tree tree,`
`TreeNodeToSend ** nodes)`

Inserts into the vector nodes the nodes to send.

Parameters

<i>tree</i>	The trees from which to get the nodes
<i>nodes</i>	The vector in which the nodes are put

2.4.5.7 tree_merge() `void tree_merge (`
`Tree * dest,`
`Tree source)`

Merge the trees dest and source and store the result in dest. The source tree is modified. It is a wrapper for.

See also

[tree_merge_dfs\(\)](#)

Parameters

<i>dest</i>	The destination tree
<i>source</i>	The source tree

2.4.5.8 tree_merge_dfs() `void tree_merge_dfs (`
 [Tree](#) * *dest*,
 [Tree](#) *source*,
 int *nd*,
 int *ns*)

Merge the subtree of *dest* rooted in node with id *nd* with the subtree of *source* rooted in *ns* and store the result in *dest*. Also the source tree is modified.

Parameters

<i>dest</i>	The destination tree
<i>source</i>	The source tree
<i>nd</i>	Id of the node in the destination tree
<i>ns</i>	Id of the node in the source tree

2.4.5.9 tree_new() `Tree tree_new ()`

Instantiate a new tree.

Returns

The new tree

2.4.5.10 tree_node_free() `void tree_node_free (`
 [TreeNode](#) * *node*)

Free a tree node.

Parameters

<i>node</i>	Pointer to the node to free
-------------	-----------------------------

2.4.5.11 tree_node_new() `TreeNode* tree_node_new (`
 `int key,`
 `int value,`
 `int parent)`

Instantiate a new node of the tree.

Parameters

<i>key</i>	The key of the node
<i>value</i>	The value of the node
<i>parent</i>	The parent of the node in the tree

Returns

Pointer to the node created

2.4.5.12 tree_print() `void tree_print (`
 `Tree tree)`

Print the tree.

Parameters

<i>tree</i>	The tree to print
-------------	-------------------

2.5 src/types.h File Reference

Definition of datatypes used by the program.

```
#include "cvector/cvector.h"  
#include "hashmap/hashmap.h"  
#include <stdbool.h>
```

Typedefs

- typedef map_t [SupportMap](#)
Map from Item to the corresponding support.
- typedef map_t [IndexMap](#)
Map from Item to its id.
- typedef uint8_t * [Item](#)
Item of a transaction.
- typedef [Item](#) * [Transaction](#)
Transaction of items.
- typedef [Transaction](#) * [TransactionsList](#)
List of transactions.

2.5.1 Detailed Description

Definition of datatypes used by the program.

2.5.2 Typedef Documentation

2.5.2.1 [IndexMap](#) typedef map_t [IndexMap](#)

Map from Item to its id.

2.5.2.2 [Item](#) typedef uint8_t* [Item](#)

Item of a transaction.

2.5.2.3 [SupportMap](#) typedef map_t [SupportMap](#)

Map from Item to the corresponding support.

2.5.2.4 Transaction `typedef Item* Transaction`

Transaction of items.

2.5.2.5 TransactionsList `typedef Transaction* TransactionsList`

List of transactions.

2.6 src/utils.h File Reference

Utility functions.

Functions

- `int min (int a, int b)`
Minimum between two integers.
- `int max (int a, int b)`
Maximum between two integers.
- `int ulength (uint8_t *s)`
Length of an Item.

2.6.1 Detailed Description

Utility functions.

2.6.2 Function Documentation

2.6.2.1 max() `int max (`
 `int a,`
 `int b)`

Maximum between two integers.

Parameters

<i>a</i>	first parameter
<i>b</i>	second parameter

Returns

maximum between a and b

2.6.2.2 min() `int min (
 int a,
 int b)`

Minimum between two integers.

Parameters

<i>a</i>	first parameter
<i>b</i>	second parameter

Returns

minimum between a and b

2.6.2.3 ulength() `int ulength (
 uint8_t * s)`

Length of an Item.

Parameters

<i>s</i>	Item
----------	------

Returns

Number of characters in s

